

A

- adjacency lists and matrices 443
- adjacent vertices 437
- adversarial sampling 288
- agents, pipeline agent implementation 409
 - building LangGraph pipeline 414
 - generate-summary agent 413
 - intent detection agent 409
 - post-query execution 413
 - query execution agent 411
 - schema extraction agent 410
 - text-to-Cypher agent 411
- AGGREGATE function 289, 291
- all-in-one NLP 113
- Amazon Textract 148
- AML (anti-money laundering) 303–304
- annotations
 - ingestion and processing 55–59
 - schema definition with 387
- answer first vs. reasoning first 384
- APOC (Awesome Procedures on Cypher library) 456–457
- apoc.meta.schema 377, 405
- APSP (all-pairs shortest path) algorithm 479
- attention mechanisms 292

B

- betweenness 247
- bottom-up approach 34
- bridge members, defined 281
- bridges, defined 441

C

- CameraEvents 423
- candidate selection (CS) 132, 188, 192
- CD (candidate disambiguation) 188, 194
 - detecting shortest paths 196
 - disambiguation 202
 - summarizing textual paths 201
 - translating paths to text 199
- chain-of-thought prompting 384
- CKG (Clinical Knowledge Graph) 91
- classification prompts 370
 - design considerations for 376
- Claude.ai 29
- clinical applications of KGs 90–94
- closeness centrality 244
- clustering 216
- CO_OFFENDS_WITH relationship 381
- communication networks 438
- community detection 216–217
- community structure 281
- complete graphs 437
- conceptual search 140, 159
- conductance, defined 78
- confusion matrices 225–227, 252–253, 303, 316, 332–333
- contextual sentences 204
- coreference resolution 5
- CRISP-DM (Cross-Industry Standard Process for Data Mining) 35
- CrossEntropyLoss function 313
- CS (candidate selection) 132, 188, 192
- cycles, defined 438
- Cypher 455
 - from text to 386–391

D

data annotation complexity and cost 113
 data understanding 141–145
 domain ontologies 142–145
 unstructured data 141
 data-centric paradigm 8, 102
 data-driven applications, paradigm shift in 10
 data-related questions 372
 DBSCAN (density-based spatial clustering of
 applications with noise) 217
 de-black-boxing 128
 deductive reasoning 28
 degree 237
 calculation 197
 of vertex 437
 density of pathway 78
 Dijkstra's algorithm 243
 directed graphs 436
 disambiguation, ingesting domain
 ontology 182–185
 disease ontology 472
 evaluating effect of normalization 476
 normalization with UMLS and
 scispaCy 473
 DL (deep learning) 273
 documentation-related questions 372
 documents, processing and ingesting 148
 domain ontology, ingesting 182–185
 DWPC (degree-weighted path count) 82, 490

E

EC (European Commission) 140
 ECDC (European Centre for Disease Prevention
 and Control) 130, 140, 181
 EDQM (European Directorate for the Quality of
 Medicines & HealthCare) 140
 egonet (ego-centered network) 236
 EHRs (electronic health records) 59, 90
 embedding generation 327
 embeddings 344
 in graph representation learning 273–278
 in knowledge graphs 285–289
 encoder-decoder architecture 310, 326
 decoder 313, 330
 encoder 311, 328
 encoder-decoder model 279–281
 decoder 280
 encoder 279
 Node2Vec 280–281
 power of framework 280
 entities, named entity recognition 100

entity co-occurrences, generating 157
 expert-emulating investigation 422–430
 context-aware request refinement 428
 historical record analysis 430
 identifying initial case 423
 spatial analysis of surveillance coverage 425
 explainability 10, 127
 explorability 128

F

fallback options 369
 false negatives 287
 feature blindness 284
 feature engineering. *See* graph feature
 engineering
 feature-based information 290
 feedback and complaints 372
 fine-tuning costs 113
 forward method 311

G

GATConv layer 313
 GATs (graph attention networks) 303
 GDL (geometric deep learning) 276
 GDS (Graph Data Science) library 74, 176, 195,
 457–458
 GenAI (generative artificial intelligence) 3
 generate-summary agent 413
 geodesic (or shortest) path 242
 geometry of embeddings 274–276
 GIS (geographic information system) 368
 GNNs (graph neural networks) 210, 273
 generalized aggregation and update
 methods 292–299
 graph processor, data preparation 305
 link prediction 304, 317–319, 326, 330–332
 message passing 289–292
 node classification 302, 304, 310, 313–314, 316,
 319
 synergy of LLMs and 299
 GO (gene ontology) 84
 graph
 classification 217, 228
 clustering 217, 229–231
 completion 215
 density 241
 graph feature engineering 233
 betweenness 247
 closeness centrality 244
 degree 237

- geodesic (or shortest) path 242
- manual relationship features 254–262
- manual node features 235–253
- PageRank 249
- prediction 251
- semiautomated feature extraction 263–271
- triangles 239
- graph processor, data preparation 305, 319
- Graph RAG 347
- graph-based entity resolution 120–122
- graphs 435
 - as models of networks 438–443
 - defined 435–438
 - representing 443–446
- Gremlin 14
- GRL (graph representation learning) 272–273
 - embeddings in 273–278
 - encoder-decoder model 279–281
 - generalized aggregation and update methods 292
 - message passing 289–292
 - shallow embeddings 283–284

H

- hallucinations 10, 102
- heterogeneous GNN model 327
- heterophily 213
- Hetionet graph 81, 84–89
- HMDD (Human miRNA Disease Database) 465
- homophily 213, 442
- HPO (Human Phenotype Ontology) 40, 141, 145
 - ingesting 155
 - repository 44

I

- i.i.d. (independent and identically distributed) 212
- IASs (intelligent advisor systems)
 - definition of 20
 - KG approach to 33–36
 - machine learning on graph to build 209
 - using domain ontologies 142–145
 - using unstructured data 141
- ICIJ (International Consortium of Investigative Journalists) 440
- in-degree 437
- inductive learning 277
- inductive reasoning 28
 - using ML 32

- information maps 439
- informational networks 438
- ingesting
 - domain ontology 182–185
 - HPO ontology 155
 - SNOMED ontology 152
- intelligent systems 17
 - categories of 20
 - characteristics of 23
 - defined 18, 20
 - designing 20–23
 - knowledge acquisition and representation 24
 - reasoning 27–30
 - reasoning engines 31–33
- intent detection 367–376
 - agent 409
 - broader spectrum of questions 372–376
 - classifying by visualization type 368–370

J

- jumping knowledge networks 298

K

- K-means 217
- KG retriever 348–349
- KG-based interpretability and discovery 140
- KGs (knowledge graphs) 4–6, 129
 - approach to intelligent agent systems 33–36
 - asking questions with natural language
 - from schema to LLM-ready context 376–383
 - intent detection 367–376
 - LLM reasoning 383–392
 - using metadata for enhanced querying 366
- building 41–46, 55–59
 - annotation ingestion and processing 55–59
 - business and domain understanding 41
 - data understanding 44–46
 - ontology ingestion and processing with neosemantics 52–55
- building data-driven applications using 12–14
- conversational AI for customer support 13
 - deciding whether to use a KG 14
 - drug discovery and development 13
- building from structured sources 462
 - disease ontology 472
 - exploring and analyzing miRNA KG 487–492

KGs (knowledge graphs) (*continued*)

- building with LLMs 101–115
 - prompt engineering examples 104–109
 - traditional NLP vs. LLMs 112–114
- chatting with AI about private data 342–354
 - chatting with KG 352–354
 - graph RAG 348–351
 - reasoning agents 351
 - retrieval-augmented generation 343–345
 - vector-based RAG limitations 345–347
- clinical applications of 90–94
- creating from ontologies 39
 - querying data 59–62
 - reasoning over knowledge graphs 62
- embeddings in 285–289
 - loss function 285–288
- enabling domain experts with 358
- four pillars of 11
- intellectual network analysis 122–126
- KG-specific annotations 390
- LLMs and 8
- machine learning on 210–231
 - clustering and community detection 216
 - graph classification 217, 228
 - graph clustering 216, 229–231
 - node classification and link prediction 211–214, 220–228
- multi-omic applications of 67–80
 - creating KGs from PPI networks and protein–disease associations 69–73
 - domain-specific analysis of PPI and disease KGs 76–80
 - high-level analysis of resulting KGs 73–76
- NED (named entity disambiguation), KG-based use cases 158
- paradigm shift in data-driven applications 10
- pharmaceutical applications of 80–89
 - deep analysis of Hetionet graph 84–89
- querying, retrieval-augmented generation for 358–362
- technologies 14, 46–52
- transforming archives to 116–122
 - creating metagraphs 119
 - graph modeling 118
 - graph-based entity resolution 120–122
 - normalization and cleansing 119
- value of in LLM era 127
- knowledge acquisition and representation 24
- knowledge cutoff 339
- knowledge extraction
 - concepts of 99
 - named entity recognition 100
 - relation extraction 101

- domain-specific knowledge extraction from unstructured data 98
- from unstructured data 97
- knowledge representation 18, 98

L

LangGraph

- building QA agents with 397
 - configuring pipeline components 401
 - expert-emulating investigation 422–430
 - Schema translation service 404
 - state management design 408
 - Streamlit application 417–422
 - system architecture overview 399
- pipeline agent implementation 409
- pipeline integration layer 415
- largest pathway component 77
- leadership roles 281
- link prediction 214–215, 220–228, 254
 - analysis and evaluation 330–332
 - encoder-decoder architecture 326
 - for movie recommendations 317
 - input data 304, 318
 - with GNNs 302
- Llama 3.1 8B, setting up model with 186
- LLMOps (LLM operations) 337
- LLMs (large language models) 3, 6, 101, 115, 235
 - building data-driven applications using 12–14
 - building KGs with 101–115
 - chatting with 339–341
 - clinical decision support analysis 93
 - domain-based NED and 136
 - graph feature engineering 261
 - knowledge graphs and 8
 - reasoning 383–392
 - answer first vs. reasoning first 384
 - from text to Cypher 386–391
 - structuring output for reliable query generation 391
 - role of in reasoning engine 33
 - synergy of GNNs and 299
 - teaching 16
 - value of KGs in era of 127
- loading ontologies 152
 - ingesting HPO 155
 - ingesting SNOMED 152
- local features 236
- log_softmax function 311, 313
- logistic regression 225
- loss function 285–288
 - negative sampling 287
- Louvain modularity algorithm 75

LPA (label propagation algorithm) 230
 LPG (Labeled Property Graph) 15, 46
 representing edge properties with 49
 vs. RDF 47

M

machine learning. *See* ML (machine learning)
 mapping ontologies 152
 ingesting HPO 155
 ingesting SNOMED 152
 MATCH pattern 388
 MEAN operator 268
 medical entities, disambiguating and
 ingesting 149
 MERGE clause 470
 message passing 289–292
 basic GNN model 291
 framework 289
 motivation and intuition 290
 with self loops 291
 metadata 348
 method 329
 using for enhanced querying 366
 metagraphs, creating 119
 metapath reduction 261
 microRNA-disease association 462
 business understanding 464
 data understanding 464
 key concepts 462
 miR2Disease 465, 470
 miRNA (microRNA) information,
 importing 482
 miRNA KG (knowledge graph) 464
 exploring and analyzing 487–492
 MISIM (miRNA similarity index tool) 486
 ML (machine learning) 97, 115, 233, 440, 464
 on knowledge graphs 210–231
 graph classification 217, 228
 graph clustering 216, 229–231
 node classification and link prediction
 211–214, 220–228
 using inductive reasoning and 32
 model-centric approach 102
 model-centric paradigm 8
 Mosquito-borne flavivirus fever 172
 multi-omic applications of KGs 67–80
 creating KGs from PPI networks and
 protein–disease associations 69–73
 domain-specific analysis of PPI and disease
 KGs 76–80
 high-level analysis of resulting KGs 73–76
 multihead attention 294–296

multirelationship decoder 288
 multisource integration 65
 multistage approach 376

N

n-ary relations 49
 narrative schema representation 379
 natural language questions 363–365
 response summarization 392–395
 NE (named entity) disambiguation
 defining schema 147
 generating entity co-occurrences 157
 KG-based interpretability and discovery 167
 processing and ingesting documents 148
 processing, loading, and mapping
 ontologies 152
 uncovering new knowledge 174
 understanding data 141–145
 NED (named entity disambiguation) 130, 180
 business and domain understanding 138–140
 candidate disambiguation 194
 candidate selection 192
 conceptual search 159
 disambiguating and ingesting medical
 entities 149
 domain-based NED and LLMs 136
 end-to-end process 187
 ingesting domain ontology 182–185
 KG-based use cases 158
 limitations of traditional systems 180
 overview of 132–136
 recognition to disambiguation 130–131
 setting up model with Ollama and Llama 3.1
 8B 186
 structured knowledge-based search 162
 negative sampling 286–287
 neighborhood attention 294
 neighborhood normalization 293–294
 neighbors 437
 Neo4j 447
 cleaning 459
 Cypher 455
 installing 450–455
 installing plugins 456–459
 overview 448–450
 neosemantics 52–55
 NER (named entity recognition) 100, 117, 129,
 187–188
 networks, graphs as models of 438–443
 networkx
 function 77
 graph 245
 library 243

neural message passing 289
 NLG (natural language generation) 13
 NLP (natural language processing) 7, 33, 97,
 180, 337, 473
 traditional 112–114
 no parameter sharing 284
 node classification 211–214, 220–228
 analysis and evaluation 313–314, 316
 encoder-decoder architecture 310
 for anti-money laundering applications 304
 graph processor, data preparation 319
 input data 304
 with GNNs 302
 node features, manual 235
 node-based combination 254–255
 node-based representation 255–256
 node-focused tasks, defined 211
 Node2Vec 223, 280
 decoding process 281
 encoding process 280
 power of framework 281
 normalization
 disease ontology 473, 476
 layers 292
 transforming archives to KGs 119
 NOT EXISTS clause 72

O

OCR (optical character recognition) 116
 Ollama, setting up model with 186
 ontologies 15
 creating knowledge graphs from 39, 59–62
 domain ontologies 142–145
 ingestion and processing with
 neosemantics 52–55
 integration 134
 processing, loading, and mapping 152, 155
 openCypher 14
 out-degree 437

P

PageRank 249–250
 parameter inefficiency 284
 path-based features 254, 256–262
 paths, defined 438–439
 pathway analysis 88
 patient experience mapping 90
 PC (path count) 83
 PDP (path-degree product) 84

pharmaceutical applications of KGs 80–89
 deep analysis of Hetionet graph 84–89
 pipeline agent implementation 409
 building LangGraph pipeline 414
 generate-summary agent 413
 intent detection agent 409
 post-query execution 413
 query execution agent 411
 schema extraction agent 410
 text-to-Cypher agent 411
 PLMs (pretrained language models) 7
 policing domain
 enabling domain experts with knowledge
 graphs 358
 querying knowledge graph in 357
 positional embeddings 276
 PPI (protein–protein interaction) networks 69,
 71
 creating KGs from 69–73
 prediction
 costs 112–113
 speed 112–113
 stability 107
 predictive model 36
 problem modeling 210
 prompt engineering 103, 111
 examples 104–109
 guidelines 109–112
 PyG (PyTorch Geometric) library 303, 307

Q

QA (question-answering) agents
 building pipeline 398
 building with LangGraph 397, 409, 417–422
 configuring pipeline components 401
 expert-emulating investigation 422–430
 future directions and enhancements 432–434
 pipeline integration layer 415
 state management design 408
 system architecture overview 399
 query execution agent 411
 query generation 391–392
 querying KGs (knowledge graphs)
 retrieval-augmented generation for 358–362
 with schema-based approach 363–365

R

RAC (Rockefeller Archive Center) 98, 115, 126
 RAG (retrieval-augmented generation) 337, 356,
 398
 challenges in production environment 341

- chatting with AI about private data 342–354
- chatting with LLMs 339–341
- for KG querying 358–362
- random walk generation 281
- RDF (Resource Description Framework) 14–15, 46
 - representing edge properties with 49, 51
 - vs. LPG 47
- RDF-star 51
- RE (relation extraction) 101, 117
- reasoning 18, 27–30
 - deductive and inductive 28
 - over knowledge graphs 62
- reasoning engines 31–33
- ReFeX (Recursive Feature eXtraction) 263–271
- relational inference 215
- relations, extraction 101
- relationship prediction 215

S

- SAGEConv layer 312
- schema extraction agent 410
- schema extraction and representation 376–383
 - enriching schemas with descriptive annotations 380
 - overview 377–380
 - practical approach to 382
- Schema translation service 404
- schema-based approach for querying KGs 363–365
- schema-related questions 372
- schemaless 36
- scispaCy 473
- scratchpad techniques 384
- security/on-premises deployment costs 112, 114
- self loops 291
- semantic consistency 384
- semifraudulent triangles 239
- semisupervised learning 214
- shallow embeddings 283–284
- shortest paths 196–197
- SIMILAR relationships 121
- sklearn library 314
- SNAP (Stanford Network Analytics Project) 70
- SNOMED (Systematized Nomenclature of Medicine) 134, 144
 - ontology, ingesting 152
- social networks 438
- softmax function 281
- SPARQL query language 14
- sparsity, defined 286
- spatial analysis 425
- stale information 10
- StandardScaler 224, 252
- state object 399, 408–409
- Streamlit application 417–422
 - LangGraph integration 420–422
 - overview of 418
- structural embeddings 276
- structural information 290
- structured data, defined 4
- structured knowledge-based search 140, 162
- structured sources
 - building knowledge graphs from 462, 464, 472, 482
 - importing known miRNA-disease connections 465
 - microRNA-disease association 462, 464
- summarization step, response 392–395
- summarizing textual paths 201
- summarizing, defined 9
- supervised algorithms 211
- supervised learning 278
- surrounding contexts 442

T

- target entity 132
- taxonomies, defined 15
- technical database schema 378
- Tesseract OCR library 116
- text-attributed graph 348
- text-paired graph 348
- text-to-Cypher agent 411
- traditional NED systems, limitations of 180
- transductive learning 277
- transfer learning 6, 101
- transformer connections 294–296
- transforming archives to KGs 116–122
 - creating metagraphs 119
 - graph modeling 118
 - graph-based entity resolution 120–122
 - normalization and cleansing 119
- translating paths to text 199
- transportation networks 438
- triangles 239–241
- TSVs (tab-separated-values) 45, 56
- Turtle (Terse RDF Triple Language) 44
- type property 119, 153
- type-constrained sampling 288

U

UMLS (Unified Medical Language System) 68,
130, 137, 142–144, 473
undirected graphs 436
unstructured data 4, 141
 extracting domain-specific knowledge
 from 97–98
unsupervised algorithms 211
unsupervised learning 278
unweighted graphs 436
update events 417, 421
UPDATE function 289, 291

V

vector creation 281

viral marketing 443
visualization, classifying by 368–370


W

W3C (World Wide Web Consortium) 14, 46
WCC (weakly connected component) 74, 230,
476
 algorithm 122
weighted graphs 436
WHERE clause 172, 175
writing styles, defined 5

Z

Zachary Karate Club 221

An AI Agent Design for KG-Powered Question Answering



Knowledge Graphs and LLMs IN ACTION

Negro • Futia • Kùs • Montagna

Forewords by Maxime Labonne, Khalifeh AlJadda

Using knowledge graphs with LLMs reduces hallucinations, enables explainable outputs, and supports better reasoning. By naturally encoding the relationships in your data, knowledge graphs help create AI systems that are more reliable and accurate, even for models that have limited domain knowledge.

Knowledge Graphs and LLMs in Action shows you how to introduce knowledge graphs constructed from structured and unstructured sources into LLM-powered applications and RAG pipelines. Real-world case studies for domain-specific applications—from healthcare to financial crime detection—illustrate how this powerful pairing works in practice. You'll especially appreciate the expert insights on knowledge representation and reasoning strategies.

What's Inside

- Design knowledge graphs for real-world needs
- Build KGs from structured and unstructured data
- Apply machine learning to enrich, complete, and analyze graphs
- Pair knowledge graphs with RAG systems

For ML and AI engineers, data scientists, and data engineers. Examples in Python.

Alessandro Negro is Chief Scientist at GraphAware and author of *Graph-Powered Machine Learning*. **Vlastimil Kùs**, **Giuseppe Futia**, and **Fabio Montagna** are seasoned ML and AI professionals specializing in Knowledge Graphs, Large Language Models, and Graph Neural Networks.

For print book owners, all digital formats are free:
<https://www.manning.com/freebook>

“Comprehensive, pragmatic, definitive.”

—Paco Nathan, Senzing

“Builds understanding both at a theoretical and a practical level.”

—Corey L. Lanum
Visualization Partners

“An excellent introduction to building KG and LLM-powered applications.”

—Dave Bechberger, Author of
Graph Databases in Action

“Comprehensive and well thought out! The authors hit it out of the park again.”

—Sujit Pal, Elsevier



**FREE
eBook**

see first page

ISBN-13: 978-1-63343-989-4



9 781633 439894